

When Should Language Models Search? A Topology-Aware Analysis of Reasoning Strategies

Anonymous ACL submission

Abstract

Search-based reasoning and chain-of-thought (CoT) prompting offer two contrasting inference-time paradigms for language models: CoT commits to a single linear trajectory, while search explores alternative partial reasoning paths before selecting an answer. In this work, we study this trade-off using A* as a controlled and optimal instantiation of heuristic-guided search. Across eight benchmarks—StrategyQA, LogicQA, HotpotQA, GSM8K, MATH, HumanEval, MBPP, and CodeContests—and across three model scales, we find that search is not a universal upgrade over CoT; instead, the two strategies solve substantially different subsets of instances. To explain this complementarity, we introduce a *Reasoning Topology Framework* based on structural features computable before generation. A derived *branching coefficient* predicts when search helps better than dataset identity, indicating that the CoT-search boundary is primarily topology-dependent. We further train a topology-aware router that selects between CoT and A* without observing generated traces. Despite this restriction, it outperforms post-hoc routers based on semantic entropy and LLM-as-critic judgments. We attribute this to a *Fluency-Correctness Asymmetry*: fluent traces are often incorrect, making trace-level adjudication unreliable. Overall, our results suggest that adaptive reasoning should route by problem structure, deploying search only when the input topology warrants its higher cost.

1 Introduction

Chain-of-thought (CoT) prompting (Wei et al., 2022) has become the standard way to elicit step-by-step reasoning from large language

models (LLMs). By generating intermediate reasoning steps before a final answer, CoT often improves performance on arithmetic, commonsense, and symbolic tasks. However, its core generation process is still strictly *linear*: the model generates a single trace $s_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_T$, and later steps are conditioned on all earlier generated contents. This design creates three potential problems. First, an early mistake can propagate through the entire trace because the model has no built-in backtracking mechanism. Second, questions that require integrating several independent pieces of evidence can be difficult to solve with a single left-to-right narrative. Third, CoT has no explicit search objective that trades off trace length and alternative hypotheses.

These limitations motivate search-based reasoning. If the model can explore multiple partial traces, then it can reconsider early commitments, compare candidate branches, and prioritize promising reasoning states. In this paper, we study a budgeted A*-search procedure over natural-language reasoning traces. The language model generates successor steps, a step cost penalizes both trace length and low-confidence steps, and a heuristic prioritizes which partial trace to expand next.

But search is not free of problems. Exploring multiple branches can also waste computation on dead-end decompositions or introduce noise on problems that a single well-chosen chain would solve directly. In our experiments, we find that A* occasionally *hurts* by overthinking simple questions (Section 3.5). The central question is therefore not “is search better?” but rather: *which structural properties of a reasoning problem predict whether branching will help?*

Concretely, we ask the following research questions: (i) Does guided search dominate

CoT, or do they trade wins? (ii) When they disagree, is the split random or systematic? (iii) Can we predict which method will win from structural features of the input alone—before generating any reasoning trace from either CoT or the search procedure? (iv) Does such a pre-generation router actually outperform post-hoc adjudication (such as an LLM based judge) that sees full reasoning traces?

Motivating example. On StrategyQA, linear CoT can pursue a locally plausible but ultimately irrelevant chain, while A* can surface the decisive evidence earlier by exploring alternative partial traces. For example (§Figure 1), on the question “Is shrimp scampi definitely free of plastic?”, CoT misses the key fact of microplastic accumulation in seafood, whereas A* reaches that premise early and answers correctly. This illustrates the central hypothesis of the paper: CoT and search are complementary, and the real challenge is deciding when to use which. This pattern persists at scale. Across all benchmarks, CoT and A* solve different subsets of examples. On StrategyQA, A* outperforms CoT (74.80% vs. 64.52%), yet the oracle reaches 84.59%. On LogicQA, CoT is slightly stronger (45.78% vs. 44.85%), but the oracle still reaches 59.29%. Together, these gaps show that the main opportunity is *strategy selection*: different questions favor different reasoning strategies.

Why A* specifically? Our goal is not to champion A* as the best search algorithm for LLMs, but to use it as a controlled diagnostic. A* separates the key ingredients of guided search—a priority frontier, a path cost, and a heuristic estimate—without entangling them with evaluator design (as in ToT), rollout policies (as in MCTS), or fixed-width pruning (as in beam search). This makes it easier to isolate *when* the act of branching itself is responsible for gains.

Contributions. Our contributions: (1) We compare CoT and A* across eight benchmarks spanning QA, math, and code generation (StrategyQA, LogicQA, HotpotQA, GSM8K, MATH, HumanEval, MBPP, CodeContests) and three model scales (0.5B, 8B, 14B), finding that neither uniformly dominates. (2) We quantify instance-level complementarity: CoT and A* solve genuinely

Motivating Example Query: “Is shrimp scampi definitely free of plastic?” **Gold: False**

✗ Chain-of-Thought Linear generation

Commits to a chronological schema, missing the biological premise.

S1. Shrimp are caught/farmed; no conclusive evidence of significant plastic.

S2. Processing does not involve contamination.

S3. Cooking does not introduce plastic.

S4. No evidence of ingredient contamination.

S5. Regulatory standards make contamination unlikely.

Ans. YES (Incorrect)

✓ A* Search Heuristic-guided

Prioritizes the key biological premise directly.

S1. Shrimp often contain microplastics in tissues.

S2. Preparation does not remove pre-existing microplastics.

Ans. FALSE (Correct)

Figure 1: Comparison of reasoning trajectories on a StrategyQA example. CoT follows a plausible but misdirected chronological schema, while A* prioritizes the key biological fact—microplastic presence in shrimp—and reaches the correct conclusion with fewer, more relevant steps.

different subsets, and the oracle gap is large enough to make routing worthwhile. (3) We introduce 13 pre-generation topology features and show that a single derived quantity—branching coefficient—predicts A* advantage better than dataset identity ($R^2 = 0.94$). (4) We show that a topology-aware router outperforms post-hoc trace evaluation, and diagnose why through a Fluency-Correctness Asymmetry analysis.

2 Methodology

We first formalize CoT and search-based reasoning under a common notation and then describe the routing mechanisms that choose be-

tween them.

2.1 Notation and Problem Setup

Let $\mathcal{Q} = (q, C, \mathcal{O})$ denote a reasoning problem, where q is the question, C is an optional context passage, and $\mathcal{O} = \{o_1, \dots, o_K\}$ is the candidate answer set for multiple-choice tasks or the output space for free-form tasks. A *reasoning trace* $\tau = (s_1, s_2, \dots, s_T)$ is an ordered sequence of natural-language reasoning steps, and a prediction function $\phi(\tau) \in \mathcal{O}$ extracts the final answer from the trace. Let $\mathcal{M}_\theta$ denote a language model with parameters θ .

2.2 Chain-of-Thought (CoT) Reasoning

In CoT prompting, the model generates a single reasoning trace autoregressively: $\tau^{\text{CoT}} = (s_1, s_2, \dots, s_T)$ and $s_t \sim \mathcal{M}_\theta(\cdot \mid q, C, \mathcal{O}, s_{1:t-1})$. Generation terminates when the model emits a terminal answer pattern, and the final prediction is $\hat{y}^{\text{CoT}} = \phi(\tau^{\text{CoT}})$. CoT explores exactly one element of the trace space with no mechanism for backtracking or branch comparison.

2.3 A* Search over Reasoning Traces

We therefore search over the space of reasoning traces using an A*-style priority rule, but with a finite node budget and a multi-goal voting procedure for tractability.

State space. The search space is a tree $\mathcal{G} = (\mathcal{S}, \mathcal{E})$ where each node $n \in \mathcal{S}$ represents a partial reasoning state: $n = (\tau_n, d_n)$, $\tau_n = (s_1, \dots, s_{d_n})$. Here, τ_n is the partial trace stored at node n and $d_n = |\tau_n|$ is its depth. The root node n_0 has $\tau_{n_0} = \emptyset$ and $d_{n_0} = 0$. An edge $(n, n') \in \mathcal{E}$ exists when $\tau_{n'} = \tau_n \oplus s$ for some reasoning step s .

Successor generation. Given a node n with state (τ_n, d_n) , we ask the language model to propose N_c candidate next steps:

$$\{(s^{(j)}, c^{(j)})\}_{j=1}^{N_c} = \text{GS}(\mathcal{M}_\theta, q, C, \mathcal{O}, \tau_n, N_c). \quad (1)$$

GS refers to the step-generation function. Each candidate consists of step text $s^{(j)}$ and a self-reported confidence $c^{(j)} \in [0.0, 1.0]$. The candidate induces a successor node n'_j with trace $\tau_{n'_j} = \tau_n \oplus s^{(j)}$ and depth $d_{n'_j} = d_n + 1$.

Goal condition. A node is a goal node if its last step contains a terminal answer pattern:

$$\text{ISGOAL}(n) = \mathbb{1}[\exists p \in \mathcal{P} : p \subset s_{d_n}], \quad (2)$$

where $\mathcal{P}$ contains answer templates such as “The answer is [A/B/C/D]” for multiple-choice tasks.

Path cost. The cumulative cost of a node trades off reasoning length and uncertainty:

$$g(n) = \sum_{t=1}^{d_n} [1 + (1 - c_t)] = d_n + \sum_{t=1}^{d_n} (1 - c_t). \quad (3)$$

Each step contributes a base cost of 1, and lower-confidence steps incur an additional penalty. This makes shorter, higher-confidence traces cheaper than longer or more uncertain ones.

2.4 Heuristic Design and Admissibility Guarantee

The heuristic estimates the remaining cost from a partial trace to a valid goal. In our implemented A* system, we use lightweight, interpretable search-control signals rather than a learned verifier. Across datasets, the heuristic combines depth and coverage: $h(n) = h_{\text{depth}}(n) + h_{\text{coverage}}(n)$ with

$$h_{\text{depth}}(n) = \omega_d \max(0, H_{\min} - d_n), \\ h_{\text{coverage}}(n) = \omega_c (1 - E(n)). \quad (4)$$

Here $H_{\min}$ is a task-level lower bound on required reasoning hops from the root, d_n is the depth of node n , $E(n) \in [0, 1]$ is the fraction of target entities or constraints already covered by the partial trace, and $\omega_d, \omega_c > 0$ are heuristic weights. Thus, h_{depth} encourages search to continue until a minimum reasoning depth is reached, while h_{coverage} favors traces that cover more answer-relevant entities or constraints. Let $c_{\min} = 1 + (1 - c_{\max})$ be the minimum possible transition cost under the path-cost definition. Since each transition has base cost 1 plus uncertainty penalty $1 - c_t$, every transition costs at least $c_{\min}$. The confidence and coverage terms are practical search-control signals, not calibrated uncertainty estimates or exact measures of evidential completeness.

Assumption 1: Minimum-hop lower bound. Any valid goal trace requires at least $H_{\min}$ reasoning steps from the root; equivalently, any non-goal node with $d_n < H_{\min}$ requires at least $H_{\min} - d_n$ further transitions.

Assumption 2: Non-goal continuation. Any non-goal node with $d_n \geq H_{\min}$ still requires at least one additional transition before reaching a goal.

Theorem 1: Conservative admissibility. Under Assumptions 1 and 2, if $\omega_d + \omega_c \leq c_{\min}$, then the depth-plus-coverage heuristic is admissible for all non-goal nodes:

$$h(n) \leq h^*(n),$$

where $h^*(n)$ is the true minimum remaining path cost from node n to a valid goal.

Proof. The formal proof is provided in Section J.

Operational instantiation: a symbolic progress scorer. The abstract heuristic in Eq. (6)–(8) is instantiated as a dataset-agnostic progress estimator over partial reasoning states. Unlike a learned verifier or LLM-based critic, this estimator does not judge the truth of a partial trace directly. Instead, it plays a critic-like search-control role by assigning lower remaining cost to states that have made sufficient structural progress: reaching the minimum reasoning depth and covering question-relevant entities. For a question q , we extract a set of content units $\mathcal{E}_q$ consisting of quoted strings, proper-name phrases, numbers, years, and non-stopword content terms. For a node n with partial trace τ_n , coverage is computed as

$$E(n) = \frac{|\{e \in \mathcal{E}_q : e \text{ appears in } \tau_n\}|}{|\mathcal{E}_q|},$$

with $E(n) = 1$ when no content units are extracted. The implemented heuristic is

$$h(n) = \omega_d \max(0, H_{\min} - d_n) + \omega_c (1 - E(n)).$$

Thus, the heuristic is an operational instantiation of the depth-coverage template: h_{depth} estimates remaining structural progress, while h_{coverage} estimates remaining content coverage. Because the same extraction and scoring

rule is applied across datasets, the heuristic is dataset-agnostic even though its minimum-hop parameter can encode a task-level lower bound. In Appendix Section A, it is described why this is dataset agnostic (with an example) and we report a human analysis which shows this critic based heuristics underestimates human-assessed progress.

2.5 Routing Strategies

Given the complementarity between CoT and A^* , we define routing functions $\mathcal{R} : \mathcal{Q} \rightarrow \{\text{CoT}, A^*\}$. We distinguish two families.

Output-based routers. They require generating traces before deciding. *Semantic entropy* (Kossen et al., 2024): routes high-entropy (under self-consistency) instances to A^* . *LLM-as-critic*: on disagreements, a critic model (DeepSeek-R1-Distill-Qwen-14B, chosen for architectural distinctness from the LLaMA-3.1-8B backbone and local executability) compares both traces:

$$\text{verdict} = \mathcal{M}_{\text{critic}}(\text{Trace}_{\text{CoT}}, \text{Trace}_{A^*}, q, C, \mathcal{O}) \quad (5)$$

Pre-generation routers. They decide before any generation. *Random forest*: trained on hand-crafted features (question length, connectives, negation). *Topology router* (hereafter *CoT- A^* router*): XGBoost on the 13 topology features, requiring no LLM inference at routing time. It is a diagnostic test of whether the CoT/ A^* boundary is predictable from problem structure.

Terminology clarification. We distinguish three uses of evaluation/scoring in this work:

- *Rule-based progress scorer*: the deterministic heuristic $h(n)$ used to prioritize partial states during A^* search (Eq. ??). Not an LLM.
- *LLM-as-Critic*: an output-based router that inspects both completed CoT and A^* traces and selects between them via pairwise comparison.
- *LLM-as-Judge*: an alternative output-based router using a different prompting formulation emphasizing faithfulness and logical validity.

2.6 Reasoning Topology Features based Routing

To characterize the structural demands of a reasoning problem *before* any model generation, we define 13 pre-generation features computed directly from the input text. These include surface length features—question length, context length, and number of context sentences; syntactic and discourse cues—logical connectives, negation markers, comparatives, and question marks; reasoning-structure features—estimated hop count from dependency links, contextual density measured as entities per sentence, distractor count, and estimated reasoning depth; and task-format features such as the number of answer options. We also define a derived *branching coefficient*:

$$BC = \frac{c_sent \times hop}{q_len} \quad (6)$$

where c_sent is the number of context sentences, hop is the estimated reasoning-hop count, and q_len is the question length. BC summarizes how much multi-hop branching structure the input affords relative to question length. These features require no model inference and are used for both topology clustering (§3.3) and pre-generation routing (§3.4).

Router evaluation protocol. The topology-aware router (referred to as the *CoT-A* router*) is evaluated using stratified 5-fold cross-validation (seed 42), performed separately within each benchmark. For each fold, the router is trained on four folds and evaluated only on the held-out fold; no evaluation instance is used to train the router that predicts it. The router receives only the 13 pre-generation features—it does not observe CoT traces, A* traces, model confidence, predicted answers, or gold answers at inference time. This protocol evaluates *within-dataset generalization*; cross-dataset transfer is examined separately via a leave-one-benchmark-out experiment (Appendix F).

3 Results

We evaluate on eight benchmarks using three models: i) LLaMA-3.1-8B, ii) Qwen-0.5B, and iii) Qwen-2.5-14B, and we use DeepSeek-R1-Distill-Qwen-14B as the LLM critic. For

Dataset	LLaMA-3.1-8B				Qwen-0.5B				Qwen-2.5-14B	
	CoT	A*	SC	ToT	CoT	A*	SC	ToT	CoT	A*
StrategyQA	64.52	74.80	71.22	<u>72.24</u>	<u>53.60</u>	53.40	54.21	52.24	76.02	75.56
LogicQA	45.78	44.85	<u>46.10</u>	46.21	24.00	20.20	<u>22.10</u>	21.19	<u>57.91</u>	58.12
HotpotQA	76.80	80.10	78.21	<u>78.45</u>	8.40	30.00	12.21	<u>22.17</u>	60.19	83.33
GSM8K	85.04	<u>86.12</u>	86.90	60.00	<u>26.84</u>	14.48	43.22	41.02	61.56	80.59
MATH	78.00	<u>79.10</u>	79.50	40.10	32.80	16.20	41.60	49.00	44.00	47.20
<i>Code Generation</i>										
HumanEval	<u>67.48</u>	64.42	65.03	71.78	33.74	20.25	26.38	<u>33.13</u>	77.91	76.07
MBPP	39.80	36.80	38.00	<u>40.40</u>	24.00	<u>20.80</u>	21.40	23.00	46.60	45.00
CodeContests	77.40	87.00	<u>99.00</u>	91.20	49.80	68.20	<u>48.80</u>	67.80	99.60	94.40

Table 1: Consolidated results across reasoning, math, and code generation benchmarks. Scores denote accuracy (%). SC denotes Self-Consistency and ToT denotes Tree-of-Thought. Within each dataset-model block, the best single-method score is bolded and the second-highest is underlined. On CodeContests at 0.5B, A* gives a +18.4 pp gain over CoT, mirroring the search advantage for small models on HotpotQA.

QA, we use StrategyQA (Geva et al., 2021) (multi-step commonsense, boolean answers), LogicQA (Liu et al., 2020) (formal logical reasoning, four-way multiple choice), and HotpotQA (Yang et al., 2018) (multi-hop QA, free-form answers). For math, we use GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021). For code generation, we use HumanEval (Chen et al., 2021) (163 function-level tasks with detailed docstrings), MBPP (Austin et al., 2021) (500 tasks with one-sentence specifications), and CodeContests (Li et al., 2022) (500 algorithmic problems). We report exact accuracies from the current experimental pipeline, and the oracle corresponds to instance-wise selection of the correct answer between CoT and A*.

3.1 Quantitative Analysis

Main results. Table 1 shows that no single reasoning strategy dominates across datasets, models, or reasoning formats. For LLaMA-3.1-8B, A* is substantially stronger than CoT on StrategyQA (+10.28 pp) and HotpotQA (+3.30 pp), while CoT has a slight edge on LogicQA (+0.93 pp). Yet the oracle remains far above the better of CoT and A*—by +9.79, +13.51, and +14.10 points on StrategyQA, LogicQA, and HotpotQA, respectively—confirming strong instance-level complementarity.

The expanded results show that this com-

plementarity is not limited to one model or dataset family. On Qwen-0.5B, A* gives a large gain on HotpotQA (30.00 vs. 8.40), but not on StrategyQA or LogicQA, suggesting that search helps small models mainly when the task structure rewards exploration. On Qwen-2.5-14B, A* is especially effective on HotpotQA (83.33 vs. 60.19) and on GSM8K (80.6% vs. 61.6%, +19.0 pp)—the largest A* advantage at any scale on math—while self-consistency is strongest on StrategyQA and LogicQA. Overall, search is most useful for multi-hop or branch-heavy inputs, whereas sampling-based or linear methods remain competitive on more compact reasoning problems. Finally, the best-router and oracle columns show that adaptive strategy selection can exceed any fixed single method, although a substantial oracle gap remains for the topology-aware routing.

The code-generation benchmarks show a similar trade-off. On CodeContests, A* outperforms CoT at both 0.5B (+18.4 pp) and 8B (+9.6 pp), with oracle gaps of 7.8 pp and 9.2 pp respectively—substantial complementarity that the topology router exploits effectively (95.0% at 8B, recovering 87% of the oracle gap). On HumanEval and MBPP, ToT is the strongest single method, but the oracle gap remains large (74.8% at 8B, 85.9% at 14B vs. best single). We present topology routing, multi-return analysis, and category breakdowns for code in Appendix L.

Router performance. We evaluate all routing strategies on LLaMA-3.1-8B to test whether a router can choose between CoT and A* more effectively than either fixed strategy. Table 4 shows that routing closes only part of the oracle gap, but the topology-aware router is consistently strongest: on the QA benchmarks, it improves over the stronger fixed baseline by +2.55 pp on StrategyQA, +3.23 pp on LogicQA, and +2.31 pp on HotpotQA, while recovering 26.1% of the oracle gap. The pattern also extends to math, where the topology router reaches 91.10% on GSM8K and 81.52% on MATH, outperforming CoT, A*, semantic entropy, LLM-as-critic, LLM-as-judge, and random-forest routing. This contrast is important: trace-based routers yield only modest gains and sometimes underperform A*

alone, whereas the topology router uses only 13 pre-generation structural features. Thus, the CoT-search decision is better predicted from problem topology than from retrospective evaluation of completed reasoning traces.

Dataset	Both correct	CoT only	A* only	Both wrong
StrategyQA	1361 (59.4%)	328 (14.3%)	248 (10.8%)	353 (15.4%)
LogicQA	180 (27.6%)	118 (18.1%)	88 (13.5%)	265 (40.7%)
HotpotQA	298 (59.6%)	86 (17.2%)	87 (17.4%)	29 (5.8%)
HumanEval (8B)	93 (57.1%)	17 (10.4%)	12 (7.4%)	41 (25.2%)
MBPP (8B)	163 (32.6%)	36 (7.2%)	21 (4.2%)	280 (56.0%)
CodeContests (8B)	341 (68.2%)	46 (9.2%)	94 (18.8%)	19 (3.8%)
CodeContests (0.5B)	210 (42.0%)	39 (7.8%)	131 (26.2%)	120 (24.0%)

Table 2: Instance-level complementarity decomposition. “CoT only” means CoT is correct and A* is wrong; “A* only” means the reverse.

Instance-level decomposition. Table 2 makes the complementarity concrete. HotpotQA is especially striking: only 5.8% of instances are missed by both strategies, while 34.6% lie in a regime where exactly one strategy is correct. StrategyQA also exhibits sizable disagreement. LogicQA is harder overall, as shown by the 40.7% “both wrong”, but even there each strategy recovers failures of the other.

3.2 Qualitative Analysis

We manually inspect representative CoT–A* disagreement cases from StrategyQA, LogicQA, and HotpotQA and group them into recurring *reasoning regimes*. The qualitative evidence supports the quantitative complementarity in Table 2: A* helps when the input induces several plausible partial trajectories, whereas CoT remains preferable when the solution is naturally linear and executable in one coherent pass. Appendix B provides the corresponding examples, organized into A* success cases, CoT success cases, A* failure modes, and human trace-quality analysis.

When A* helps. A* is most useful when the reasoning topology is *branching*: the model must compare, preserve, or eliminate multiple partial hypotheses before committing. Appendix B shows three such regimes. In *non-linear evidence aggregation* (Appendix B.1), decisive facts do not form a single fluent chain; CoT may follow a locally plausible narrative, while search keeps alternatives open until a globally better path appears. In *hypothesis elimination*, especially multiple-choice Log-

icQA, the task is less to derive one proof than to rule out plausible distractors; search can distribute budget across options, whereas CoT may commit to the first surface-compatible choice. In *multi-hop composition*, especially HotpotQA, search preserves competing entity or relation chains until the bridge entity is found, improving evidence traversal rather than adding new knowledge.

Why these gains arise. The common mechanism is *deferred commitment*. CoT commits to a single left-to-right trajectory, which is efficient when the early steps are correct but brittle when they are only partially informative or misleading. A* treats intermediate thoughts as expandable states, allowing the model to back away from an initially promising but suboptimal path. Its advantage is therefore not generic extra computation, but better allocation of reasoning effort across competing partial solutions, as illustrated by the A* success cases in Appendix B.1.

When CoT remains preferable. Search hurts when the problem does not require exploration. Appendix B.2 shows two such regimes. In *direct factual reasoning*, the answer follows from a small set of stable facts or a short deterministic comparison; a single fluent chain is sufficient, while A* may introduce irrelevant distinctions. In *globally coherent argumentation*, especially some LogicQA cases, the solution requires maintaining one abstract frame across the whole argument. CoT can preserve this global frame in a continuous explanation, whereas A* may fragment it into locally plausible but globally weaker steps.

Characteristic failure modes of A*. The A* errors are systematic rather than random, and Appendix B.3 gives representative cases. The dominant failure is *overthinking simple problems*: when direct recall or a short inference suffices, search may add unnecessary distinctions and drift from the obvious answer. A second failure is *spurious associative expansion*, where a weak semantic link is elaborated despite limited causal relevance. A third is *budget waste through repetition or near-duplicate expansions*, where search revisits the same intermediate fact instead of exploring genuinely new possibilities. These cases show

Cluster	Profile	BC	A*	CoT	Interpretation
C1	Linear, short	1.28	81.7	82.1	CoT sufficient
C2	Mid-branch	4.24	73.5	66.2	selection critical
C3	Complex, noisy	4.02	52.3	49.1	both often fail
C4	Deep-branch	6.49	69.8	58.4	A* largest gain

Table 3: Reasoning topology clusters. BC = branching coefficient. A* advantage grows monotonically with branching complexity.




Figure 2: A* advantage (pp) vs. branching coefficient across topology clusters. Linear fit: $R^2 = 0.94$.

that search changes the *shape* of reasoning, not only its computational cost; when branching is unnecessary, this changed shape can reduce correctness.

Takeaway. The CoT-A* tradeoff is topological: A* helps when the problem structure warrants exploration, while CoT is stronger when direct, coherent commitment is enough.

3.3 Explaining Complementarity Through Topology

We formalize the qualitative intuition using the topology features from Section 2.6. Applying k -means ($k=4$) to the 13 features across all QA and math instances yields four clusters (Table 3). A* advantage increases from -0.4 pp (C1) to $+11.4$ pp (C4), with $R^2 = 0.94$ on these clusters (Figure 2). Each cluster contains instances from multiple benchmarks, indicating the CoT/A* boundary is explained by structure, not dataset identity.

Prescriptive decision rule. The cluster analysis yields a practical decision rule: *C1 (linear)*: use CoT; search overhead is unnecessary. *C2 (mid-branch)*: route adaptively; this is where router quality matters most. *C3 (complex, noisy)*: neither method is reliable; retrieval augmentation or stronger models may be needed. *C4 (deep-branch)*: use A*; search has the clearest advantage.

Router	Input	Str.QA	Log.QA	Hot.QA	CC	HE	MATH	ImpHE	MBQA	CC	QA	Gap%	Tot.QA	GSM8K	M
A* only	traces	74.80	44.85	80.10	86.12	79.10	64.42	36.80	87.00	−100	Qwen-2.5-14B				
CoT only	traces	64.52	45.78	76.80	85.04	78.00	67.48	39.80	77.40	−24.6			82.50	80.59	
Sem. Entropy	traces	74.37	47.47	81.30	86.67	78.24	65.64	37.08	88.20	−3.95			56.20	61.56	
LLM-as-Critic	traces	74.19	47.16	80.21	87.28	79.13	71.17	40.20	97.80	9.6			68.40	73.31	
LLM-as-Judge	traces	73.28	47.21	80.87	85.63	78.23	71.78	40.40	97.80	6.6			82.50	70.13	
Random Forest	features	76.25	48.20	81.25	81.43	80.12	67.48	39.38	77.30	−3.3			57.55	69.07	
Topology Router	13 feat.	77.35	49.01	82.41	91.10	81.52	72.50	40.63	95.00	40.0			75.50	61.56	
Oracle	–	84.59	59.29	94.20	Topology Router	13 feat.	74.54	58.53	96.58	100			83.17	80.82	

Table 4: Router comparison on LLaMA-3.1-8B. HE = HumanEval, CC = CodeContests. Scores denote accuracy (%). On CodeContests, the topology router achieves 95.0% (vs. A*-only 87.0%, oracle 96.2%)—recovering 87% of the oracle gap, the strongest single-dataset result. Gap% to be recomputed with corrected CC data; SE/Critic/Judge on CC pending re-run. All QA router improvements over CoT: McNemar’s $p < 0.005$. LLM-as-Judge and LLM-as-Critic prompts in Appendix Figure 3.

Topic modelling validates topology.

As independent validation, we apply Non-negative Matrix Factorisation (NMF, $k=8$) on TF-IDF representations of all 3,660 paired evaluation questions. A* dominates in 6 of 8 topic categories covering 66.2% of the corpus—these share multi-hop entity chains (entertainment, historical events, geography, biography, politics). CoT-dominant topics (33.8%) correspond to single-step commonsense categories (everyday reasoning, hypothetical scenarios). Full per-topic results with examples in Appendix H.

3.4 Topology-Aware Routing

Table 4 compares all routers and reports oracle gap recovery.

Why pre-generation routing outperforms post-hoc routing. The topology router observes no generated traces, yet outperforms routers that inspect both full reasoning traces. This reverses the intuitive expectation that seeing full traces should be more informative. The bottleneck is often not judging which completed trace looks better, but recognizing before generation which computation the problem requires. The same holds for code: on HumanEval (8B), the topology router reaches 72.5% vs. 71.8% for LLM-as-Judge, and on CodeContests (0.5B) it recovers 89.7% of the oracle gap—both without seeing any generated code (Table 5, Appendix L).

Fluency-Correctness Asymmetry. We identify a fundamental limitation of output-

Oracle	–	81.22	68.51	86.40	89.69
Qwen-0.5B					
A* only	traces	59.09	20.20	30.00	14.48
CoT only	traces	56.36	24.00	7.50	26.84
Sem. Entropy	traces	56.36	24.00	12.50	26.84
LLM-as-Critic	traces	59.09	24.00	30.00	27.90
LLM-as-Judge	traces	58.18	32.40	32.50	27.37
Random Forest	features	61.82	21.20	25.00	26.84
Topology Router	13 feat.	64.74	24.98	30.00	26.84
Oracle	–	81.82	34.20	35.00	34.04

Table 5: Router comparison at 14B and 0.5B scales. HE = HumanEval, CC = CodeContests. The topology router is highest on most datasets at both scales; at 0.5B it recovers 89.7% of the CodeContests oracle gap using only problem text. – indicates data not yet available.

based routing: in 38% of disagreement cases, CoT produces the *more fluent* trace but the *wrong* answer (Table 7). A critic that rewards coherence can therefore prefer the wrong computation. This asymmetry explains why the topology router (which sidesteps trace evaluation entirely) outperforms the LLM-as-critic.

Feature ablation: reasoning structure vs. surface statistics. To verify that the topology router captures genuine reasoning structure rather than surface artifacts, we evaluate feature subsets (Table 6). Strikingly, using only the 9 reasoning-structure features (connectives, negation, hops, comparisons, linearity, depth, branching complexity, BC)—*excluding* all length and format features—**improves** router performance from 71.8% to 75.0% CV accuracy (Gap: 41.9% $\rightarrow$ 48.5%). This demonstrates that the topology router’s predictive power comes from reasoning-structural properties, not from length or format statistics that might encode dataset identity.

3.5 Sensitivity, Compute, and Failure Analysis

Sensitivity analysis. On a stratified 300-question HotpotQA sample (seed=42), we sweep the node budget $B \in \{5, 10, 15, 20, 30\}$

Feature Set	CV Accuracy	Gap
All 14 features	0.718	41.8%
Reasoning only (9 features)	0.750	48.5%
Format only (5 features)	0.713	40.8%
No length features (12)	0.717	41.8%
BC only (1 feature)	0.709	40.0%
Top-5 reasoning features	0.745	47.4%

Table 6: Feature ablation on pooled 8B data (6 benchmarks, 2660 instances). Reasoning-only features (connectives, negation, hops, comparisons, depth, BC) outperform the full set, confirming that the router learns reasoning structure, not surface statistics.

Metric	A*	CoT
Average fluency score	0.725	0.673
Cases where more fluent	62.0%	38.0%
More-fluent but wrong	38.0% of disagreements	

Table 7: Fluency-Correctness Asymmetry on disagreement cases. Trace fluency does not predict correctness.

and branching factor $K \in \{1, 2, 3, 4\}$ (full tables in Appendix D). Accuracy improves from $B=5$ (73.0%) to $B=20$ (78.7%), then drops at $B=30$ (76.0%), indicating diminishing returns. The default $B=15$ (78.0%) lies within 0.7 pp of the peak. Even $K=1$ (no real branching) achieves 77.7%, confirming graceful degradation and exceeding self-consistency.

Compute-accuracy tradeoff. A* at the default ($K=3$, $B=15$) uses $\sim 12,452$ measured tokens/question ($\sim 24 \times$ CoT’s ~ 512). A budget-reduced variant ($K=1$) achieves competitive accuracy at 3,372 tokens/question ($\sim 7 \times$ CoT). A* is substantially more expensive; its use is justified only when topology predicts branching is useful, motivating pre-routing as a compute-saving mechanism (full table in Appendix E).

Failure analysis. Among A* failures on instances CoT solves correctly, 91.1% are *overthinking* (search continues past correct evidence, accumulating distractors) and 8.9% are *premature commitment* (search locks onto an early branch). CoT fails by committing too early; A* fails by continuing too long. The same branching mechanism that helps on multi-hop problems hurts on linear ones (examples in Appendix B.3).

Human evaluation. We conduct a human evaluation of reasoning-trace quality us-

ing the criteria in §B.4. Two post-graduate annotators with experience in LLM reasoning search each labeled 100 instances from StrategyQA and LogicQA. Across datasets and metrics, Cohen’s κ ranges from 0.59 to 0.68, indicating fair agreement. Two annotators labeled 100 disagreement instances per dataset on *coverage* (1–3) and *relevance* (1–3), with Cohen’s κ : 0.59–0.68. The correct method consistently achieves higher coverage (e.g., StrategyQA: CoT-correct 3.00 vs. A* 1.20; A*-correct 2.71 vs. CoT 1.63), confirming that disagreements are driven by coverage of the required reasoning path rather than surface quality (full results in Appendix B.4).

4 Related Work

CoT and structured reasoning. Chain-of-thought prompting (Wei et al., 2022) reasons through a single linear trace, while self-consistency (Wang et al., 2023), Tree-of-Thoughts (Yao et al., 2024), and Graph-of-Thoughts (Besta et al., 2024) add sampling, tree search, or graph-structured deliberation. Prior work mostly asks whether richer reasoning structures improve aggregate accuracy, and recent work shows that CoT itself is not uniformly beneficial (Meincke et al., 2025). We ask the complementary question: when does moving from linear reasoning to search help, and when does it hurt?

Search-augmented reasoning and inference-time compute. MCTS-style reasoning (Hao et al., 2023), beam search (Xie et al., 2024), LATS (Zhou et al., 2024), MCTSr (Zhang et al., 2024), and rStarMath (Guan et al., 2025) demonstrate the value of explicit exploration. (Snell et al., 2024) show that scaling test-time compute can be more effective than scaling parameters, but leave open the question of *which* problems benefit from additional compute. Self-Refine (Madaan et al., 2023) and reasoning-focused models like DeepSeek-R1 (DeepSeek-AI, 2025) allocate variable compute adaptively, but do not provide pre-generation prediction of when this helps. We use A* as a controlled diagnostic and show that search gains are topology-dependent: branching, multi-hop, and hypothesis-elimination problems benefit

Method	Structure	Front.	Heur.	Pre-rte	Limitation
CoT	linear	✗	✗	✗	premature commit
SC	multi-linear	✗	✗	✗	unguided sampling
ToT	tree/BFS	✓	eval.	✗	evaluator entangled
MCTS	tree+rollout	✓	value	✗	many choices
A* (ours)	priority	✓	✓	✓	higher compute

Table 8: Methodological comparison of reasoning strategies. Front. = explicit frontier; Heur. = heuristic guidance; Pre-rte = pre-generation routing studied. A* isolates the role of heuristic-guided search without entangling evaluator design or rollout policy.

most, while simpler linear cases can degrade.

Over-exploration and search failures.

Work on overthinking and excessive reasoning often treats over-exploration as an efficiency issue (Sui et al., 2025; Wang et al., 2025b; Ma et al., 2025), and (Wang et al., 2025a) identify over- and under-exploration as dual pitfalls of tree search. We show the problem is also one of correctness: 91.1% of A* failures on CoT-correct instances arise from overthinking, where search continues beyond sufficient evidence, accumulates distractors, or revisits near-duplicate states.

Routing and critic limitations. Semantic entropy (Kossen et al., 2024), LLM-as-judge methods (Zheng et al., 2023), and process supervision (Lightman et al., 2023) rely on generated traces to assess reasoning quality, while mixture-of-experts models (Shazeer et al., 2017) route between neural modules. We instead route between *reasoning algorithms* before generation, using input topology (Table 8). This is crucial because our Fluency–Correctness Asymmetry shows that fluent traces can be wrong: in 38% of disagreement cases, the more fluent trace gives the incorrect answer. Our contribution is a diagnostic framework for deciding *when* search should be used.

5 Conclusion

Across eight benchmarks—spanning QA, math, and code generation—and three model scales, we find that CoT and A* are complementary rather than hierarchical: the gap between the oracle and the best single method

is consistently large. The key explanatory variable is branching coefficient, not dataset identity—a structural finding that holds across all three domains. Perhaps most surprisingly, a router that sees only the raw input (no traces) outperforms critics that inspect both full reasoning chains, because trace fluency misleads post-hoc judges in a systematic way. Code generation exhibits the same pattern: on CodeContests at 8B, the topology router reaches 95.0% accuracy (vs. A*-only 87.0%), recovering 87% of the oracle gap—the strongest single-dataset result. At 0.5B, A* gives +18.4 pp on CodeContests, and the overall 8B gap recovery across all eight benchmarks reaches 40.0%. These results suggest that the next step for adaptive inference is not better trace evaluation but better problem-structure recognition. Concrete next steps include learned topology embeddings and code-specific structural features (e.g., AST complexity, specification ambiguity).

Limitations

Our findings are conditioned on a single budgeted A* implementation; a different search policy (MCTS, beam search) could shift where the complementarity boundary lies. The 13 topology features are hand-designed proxies—interpretable but unlikely to capture the full latent structure of reasoning difficulty. They also rely on English-language parsing tools, limiting direct transfer to other languages or code. The critic-based heuristic depends on DeepSeek-R1-Distill-Qwen-14B; although the topology router reduces this dependence, it does not eliminate it entirely. On the cost side, A* at default settings uses $\sim 24\times$ the tokens of CoT, and our token accounting covers only a 300-question ablation sample. Domain coverage spans QA, math, and code generation; long-form generation remains untested. On code tasks, the NL-oriented topology features work well on HumanEval (rich specifications) but degenerate on MBPP (ultra-short prompts); code-specific topology features (e.g., AST complexity, edge-case count) remain future work. Finally, the topology router is an XGBoost classifier on fixed features; a learned router that jointly optimizes routing and search would be a natural next step.

References

- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. 2021. [Program synthesis with large language models](#). *arXiv preprint arXiv:2108.07732*.
- Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Niewiadomski, Piotr Nyczyk, and 1 others. 2024. Graph of thoughts: Solving elaborate problems with large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, and 1 others. 2021. [Evaluating large language models trained on code](#). *arXiv preprint arXiv:2107.03374*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. [Training verifiers to solve math word problems](#). *Preprint*, arXiv:2110.14168.
- DeepSeek-AI. 2025. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*.
- Mor Geva, Daniel Khashabi, Elad Segal, Tushar Khot, Dan Roth, and Jonathan Berant. 2021. Did aristotle use a laptop? a question answering benchmark with implicit reasoning strategies. In *Transactions of the Association for Computational Linguistics*, volume 9, pages 346–361.
- Xinyu Guan, Li Li, Yongle Chen, Shuo Shao, Liu Gong, Zhenyu Wang, Zhongxin Sun, Shuai Zhai, Weiwei Ai, and Weizhi Li. 2025. rstar-math: Small LLMs can master math reasoning with self-evolved deep thinking. *arXiv preprint arXiv:2501.04519*.
- Shibo Hao, Yi Gu, Haodi Ma, Joshua Hong, Zhen Wang, Daisy Wang, and Zhiting Hu. 2023. Reasoning with language model is planning with world model. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. [Measuring mathematical problem solving with the math dataset](#). *Preprint*, arXiv:2103.03874.
- Jannik Kossen, Jiatong Han, Muhammed Razzak, Lisa Schut, Shreshth Malik, and Yarin Gal. 2024. [Semantic entropy probes: Robust and cheap hallucination detection in llms](#). *Preprint*, arXiv:2406.15927.
- Daniel D Lee and H Sebastian Seung. 1999. Learning the parts of objects by non-negative matrix factorization. *Nature*, 401(6755):788–791.
- Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, and 1 others. 2022. Competition-level code generation with AlphaCode. *Science*, 378(6624):1092–1097.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2023. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*.
- Jian Liu, Leyang Cui, Hanmeng Liu, Dandan Huang, Yile Wang, and Yue Zhang. 2020. LogiQA: A challenge dataset for machine reading comprehension with logical reasoning. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence*.
- Xinyin Ma, Guangneng Jiang, Yueqian Gong, and 1 others. 2025. CoT-Valve: Length-compressible chain-of-thought tuning. *arXiv preprint arXiv:2502.09601*.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, and 1 others. 2023. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*.
- Lukas Meincke, Simon Razniewski, and 1 others. 2025. Chain-of-thought is not always beneficial: When and why CoT helps or hurts. *arXiv preprint arXiv:2502.03601*.
- Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. 2017. [Outrageously large neural networks: The sparsely-gated mixture-of-experts layer](#). *Preprint*, arXiv:1701.06538.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*.
- Yang Sui, Xiao Yu, Jiawei Zhang, and 1 others. 2025. Stop overthinking: A survey on efficient reasoning for large language models. *arXiv preprint arXiv:2503.16419*.
- Ante Wang, Linfeng Song, Ye Tian, Dian Yu, Haitao Mi, Xiangyu Duan, Zhaopeng Tu, Jin-song Su, and Dong Yu. 2025a. [Don’t get lost in](#)

the trees: Streamlining LLM reasoning by overcoming tree search exploration pitfalls. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 23946–23959, Vienna, Austria. Association for Computational Linguistics.

Jinhao Wang, Dongming Yang, and 1 others. 2025b. Fetch: A memory-efficient exploration method for tree search in LLM reasoning. *arXiv preprint arXiv:2502.12345*.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V Le, Ed H Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations*.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, and 1 others. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*.

Yuxi Xie, Kenji Kawaguchi, Yiran Zhao, James Xu, Min-Yen Kan, Junxian He, and Michael Xie. 2024. Self-evaluation guided beam search for reasoning. In *Advances in Neural Information Processing Systems*.

Zhiping Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D Manning. 2018. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Tom Griffiths, Yuan Cao, and Karthik Narasimhan. 2024. Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems*.

Di Zhang, Jiatong Wu, Jingdi Lei, and 1 others. 2024. Accessing GPT-4 level mathematical olympiad solutions via Monte Carlo Tree Self-Refine with LLaMa-3 8B. *arXiv preprint arXiv:2406.07394*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, and 1 others. 2023. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*.

Andy Zhou, Kai Yan, Michal Shlapentokh-Rothman, Haohan Wang, and Yu-Xiong Wang. 2024. LATS: Language agent tree search unifies reasoning, acting, and planning in language

models. *Proceedings of the International Conference on Machine Learning*.

A More on admissibility and generalizability of the heuristics

Admissibility and consistency. Under the minimum-hop and non-goal-continuation assumptions, this heuristic is admissible when its weights are bounded by the minimum transition cost. In our implementation, $H_{\min} = 2$, $\omega_d = \omega_c = 0.3$, and $c_{\max} = 0.99$, so the minimum transition cost is

$$c_{\min} = 1 + (1 - c_{\max}) = 1.01.$$

The largest possible heuristic value at the root is

$$2\omega_d + \omega_c = 0.9,$$

which is below the minimum two-step path cost $2c_{\min} = 2.02$. After one step, the heuristic is at most $\omega_d + \omega_c = 0.6 < c_{\min}$, and after the minimum-hop threshold the coverage term is at most $\omega_c = 0.3 < c_{\min}$. Hence the heuristic lower-bounds the remaining path cost under the stated assumptions and is admissible. It is also consistent, since a single transition can reduce the heuristic by at most $\omega_d + \omega_c = 0.6$, which is smaller than the minimum transition cost $c_{\min} = 1.01$.

Human validation of heuristic progress.

We additionally validate whether the symbolic progress heuristic behaves conservatively on intermediate reasoning states. We sample 500 partial states produced during A* search and ask human annotators to assign a normalized progress score indicating how close the partial trace is to a valid answer. The model used is LLaMA-3.1-8B and the datasets are 100 StrategyQA, 100 HotpotQA, and 300 LogicQA. We then compare this human-assessed progress with the heuristic’s implied progress score. The heuristic underestimates progress on average (0.43 vs. 0.58, paired t -test $p < 0.001$), with an overestimation rate of 12.3% and maximum observed overestimation $\epsilon_{\max} = 0.14$. Thus, the heuristic is conservative in expectation: it tends not to overstate the completeness of a partial trace. This supports its use as a safe search-control signal rather than as an unconstrained learned verifier.

Algorithm 1 Dataset-Agnostic Symbolic Progress Heuristic

Input: Question q ; partial reasoning state $n = (\tau_n, d_n)$; minimum-hop lower bound $H_{\min}$; weights ω_d, ω_c

Output: Heuristic value $h(n)$

Step 1: Extract question content units.

Initialize $\mathcal{E}_q \leftarrow \emptyset$. Add all quoted strings in q to $\mathcal{E}_q$. Add all proper-name phrases in q to $\mathcal{E}_q$. Add all numbers and years in q to $\mathcal{E}_q$. Add all non-stopword content terms of length greater than three to $\mathcal{E}_q$.

Step 2: Compute entity/content coverage. if $\mathcal{E}_q = \emptyset$ then

| $E(n) \leftarrow 1$

else

| $C(n) \leftarrow \{e \in \mathcal{E}_q : e \text{ appears in } \tau_n\}$
 | $E(n) \leftarrow \frac{|C(n)|}{|\mathcal{E}_q|}$

Step 3: Compute remaining structural progress. $r_{\text{depth}}(n) \leftarrow \max(0, H_{\min} - d_n)$

$h_{\text{depth}}(n) \leftarrow \omega_d r_{\text{depth}}(n)$

Step 4: Compute remaining content progress. $h_{\text{coverage}}(n) \leftarrow \omega_c(1 - E(n))$

Step 5: Return total heuristic. $h(n) \leftarrow h_{\text{depth}}(n) + h_{\text{coverage}}(n)$ **return** $h(n)$

B Qualitative Examples

This appendix provides the detailed examples corresponding to the qualitative conclusions summarized in Section 3.2.

B.1 Patterns of A* Success and CoT Failure

We identify three recurring settings where A* succeeds while CoT fails.

Group 1: Non-linear evidence aggregation. These are questions where the answer requires combining several facts that do not naturally form a single left-to-right narrative.

StrategyQA — Non-linear Evidence Aggregation

Q: *Is shrimp scampi definitely free of plastic?* (Gold: No)

CoT ($\times$ predicts YES): Proceeds through five sequential steps analyzing

harvesting, processing, cooking, ingredients, and regulations. Each step independently concludes no significant plastic risk. The linear structure prevents cross-referencing marine biology evidence with the food preparation chain.

A* ($\checkmark$ predicts No): Step 1 directly identifies microplastic bioaccumulation in shrimp tissue. Step 2 notes that cooking does not remove microplastics. The heuristic guides the search toward the most relevant evidence node, avoiding the exhaustive but misleading step-by-step analysis.

LogicQA — Non-linear Evidence Aggregation

Q: *Based on the above statement [spatial arrangement of communities around Civic Park], which of the following can be derived?* (Gold: A)

CoT ($\times$ predicts D): Analyzes spatial relationships sequentially, accumulating errors in relative positioning. By the fourth step, the accumulated imprecision leads to selecting an incorrect spatial relationship.

A* ($\checkmark$ predicts A): Explores multiple spatial constraint combinations in parallel. The search finds a consistent assignment by simultaneously satisfying the “southwest of” and “southeast of” constraints, correctly deriving that Civic Park is north of the administrative service area.

Group 2: Hypothesis elimination. A* is also effective when the problem naturally decomposes into evaluating or eliminating alternatives.

LogicQA — Hypothesis Elimination

Q: *Which of the following is most similar to the argument: “All landscape rooms can see the landscape, but Li Wenbing’s family can’t see the landscape. Therefore, Li Wenbing’s family is not*

1036 a landscape room.” (Gold: D)

CoT (× predicts C): Identifies the logical structure (modus tollens) but then incorrectly maps it to option C (a modus ponens structure) due to surface similarity.

A* (✓ predicts D): Separately evaluates options through different branches. One branch identifies the negative inference pattern in the premise and correctly maps it to option D (“Anyone who meets conditions can apply; Sun Wen did not apply; therefore Sun Wen did not meet conditions”), which preserves the modus tollens structure.

1037 **Group 3: Multi-hop question answering.**
1038 HotpotQA highlights the value of search when
1039 the answer depends on chaining facts across
1040 sources.

HotpotQA — Multi-hop Chaining

Q: *What government position was held by the woman who portrayed Corliss Archer in the film Kiss and Tell?* (Gold: Chief of Protocol)

CoT (× predicts “Ambassador”): Correctly identifies Shirley Temple as the actress but then retrieves an incorrect government position from parametric memory, confusing ambassadorial roles with her protocol role.

A*: Explores multiple evidence paths for Shirley Temple’s government career. While the trace is imperfect, the tree search structure allows it to consider multiple candidate positions before committing.

1041 B.2 Patterns of CoT Success and A* Failure

1042 We identify two common settings where CoT
1043 remains better than search.

1046 **Group 1: Direct factual reasoning.**
1047 When the answer follows directly from familiar facts, search can introduce unnecessary
1048 branching and semantic drift.
1049

StrategyQA — Direct Factual Reasoning

Q: *Is a Boeing 737 cost covered by Wonder Woman (2017 film) box office receipts?* (Gold: YES)

CoT (✓ predicts YES): Directly retrieves the production budget (~\$150M), box office (~\$822M), and Boeing 737 cost (~\$100M), concluding that the receipts exceed the aircraft cost.

A* (× predicts NO): The first search step focuses on whether box office receipts are “associated with” aircraft purchases, misreading the question as one about financial transactions rather than magnitude comparison.

StrategyQA — Direct Factual Reasoning

Q: *Is average number of peas in a pod enough commas for a billion?* (Gold: YES)

CoT (✓ predicts YES): Correctly identifies that a billion (1,000,000,000) requires three commas and that a pea pod usually contains 7–9 peas, so the answer is yes.

A* (× predicts NO): Incorrectly states that a billion has 12 zeros and becomes confused by the scale comparison.

1051 **Group 2: Fluent world knowledge and argumentation.** When the problem is best
1052 solved by a smooth, coherent narrative, CoT
1053 often remains more reliable than fragmented
1054 search steps.
1055
1056

LogicQA — Fluent Argumentation

Q: *Which of the following is the best argument against the claim that rich Chinese are transferring property abroad?* (Gold: B)

CoT ($\checkmark$ predicts B): Systematically compares the answer options and identifies the strongest counterargument.

A* ($\times$ predicts A): The first search step identifies option A as relevant, and the heuristic keeps the search near that branch without sufficiently comparing the full option set.

B.3 A* Failure Modes on Simple Cases

We observe three recurrent failure modes for A* on examples that CoT solves easily.

Failure mode 1: Overthinking simple questions.

StrategyQA — Overthinking

Q: *Does Disney have an ice princess?* (Gold: YES)

A* ($\times$ predicts No): Step 1 correctly identifies Elsa from *Frozen* as having ice powers. Step 2 notes that she is called “The Ice Queen” in promotional material. The search then turns the easy recall problem into an unnecessary semantic distinction between “queen” and “princess,” yielding the wrong answer.

Failure mode 2: Spurious correlations in search.

StrategyQA — Spurious Correlation

Q: *Is Menthol associated with Thanksgiving?* (Gold: No)

A* ($\times$ predicts YES): Step 1 associates menthol with cough drops used in cold weather. Step 2 associates Thanksgiving with cold weather in North America. Step 3 turns these weak associations into a spurious chain, effectively reasoning menthol $\rightarrow$ cold weather $\rightarrow$ Thanksgiving.

Failure mode 3: Step repetition.

Algorithm 2 Budgeted A* Search over Reasoning Traces

Require: Problem $\mathcal{Q} = (q, C, \mathcal{O})$, model $\mathcal{M}_\theta$, depth limit $D_{\max}$, node budget B , candidates N_c

- 1: Initialize root $n_0 \leftarrow (\emptyset, 0)$, $g(n_0) \leftarrow 0$, $h(n_0) \leftarrow h(n_0)$ from the heuristic (Section 2.4)
- 2: OPEN $\leftarrow \{n_0\}$ (priority queue on $f(n) = g(n) + h(n)$)
- 3: VISITED $\leftarrow \emptyset$, GOALS $\leftarrow \emptyset$, explored $\leftarrow 0$
- 4: **while** OPEN $\neq \emptyset$ and explored $< B$ **do**
- 5: $n \leftarrow \text{OPEN.pop_min}()$
- 6: **if** sig(n) $\in$ VISITED **then**
- 7: continue
- 8: **end if**
- 9: VISITED $\leftarrow \text{VISITED} \cup \{\text{sig}(n)\}$
- 10: explored $\leftarrow \text{explored} + 1$
- 11: **if** ISGOAL(n) **then**
- 12: GOALS $\leftarrow \text{GOALS} \cup \{n\}$
- 13: continue
- 14: **end if**
- 15: **if** $d_n \geq D_{\max}$ **then**
- 16: continue
- 17: **end if**
- 18: $\{(s^{(j)}, c^{(j)})\}_{j=1}^{N_c} \leftarrow \text{GENERATE_STEPS}(\mathcal{M}_\theta, q, C, \mathcal{O}, \tau_n, N_c)$
- 19: **for each** $(s^{(j)}, c^{(j)})$ **do**
- 20: $n' \leftarrow (\tau_n \oplus s^{(j)}, d_n + 1)$
- 21: $g(n') \leftarrow g(n) + 1 + (1 - c^{(j)})$
- 22: OPEN.push($n', g(n') + h(n')$)
- 23: **end for**
- 24: **end while**
- 25: **return** weighted vote over GOALS

StrategyQA — Step Repetition

Q: *Will the Albany in Georgia reach a hundred thousand occupants before the one in New York?* (Gold: No)

A* ($\times$ predicts YES): Step 2 states the population of Albany, New York ($\sim 98,000$). Step 3 largely repeats the same fact. The repetition wastes node budget that should have been used to compare growth trajectories.

B.4 Human evaluation

We conduct a human evaluation of reasoning-trace quality using the criteria in §B.4. Two post-graduate annotators with experience in LLM reasoning research each labeled 100 instances from StrategyQA and LogicQA (Table 9). Across datasets and metrics, Cohen’s κ ranges from 0.59 to 0.68, indicating fair agreement.

Annotation criteria. *Coverage* measures whether a trace reaches the required inferential path and supports the final conclusion. We use a 3-point scale: 1 if the reasoning is mostly incorrect and does not reach the conclusion, 2 if it is mostly correct but still fails to reach the conclusion, and 3 if both the reasoning and conclusion are correct.

Relevance measures how focused the trace is on concepts actually needed to answer the question. We score it coarsely on a 3-point scale based on the proportion of relevant steps, where relevant steps remain on-topic and avoid introducing unnecessary concepts.

Relevance is less diagnostic. In StrategyQA, CoT remains more relevant than A* when both are wrong (2.82 vs. 1.55), while in LogicQA both methods stay moderately relevant even when incorrect (2.76 for CoT, 2.53 for A*). Thus, staying on-topic is not enough: a trace may remain relevant yet still miss the decisive inferential steps. Overall, the annotations indicate that the main difference between CoT and A* on disagreement cases lies in *which method better covers the decisive reasoning steps*, while relevance is only weakly associated with correctness.

Interpretation. Refer to Table 10. Across both datasets, the method that is correct on a disagreement case consistently receives a much higher *coverage* score than the method that is wrong. On StrategyQA, when CoT is correct and A* is wrong, CoT achieves average coverage 3.00 versus 1.20 for A*; when A* is correct and CoT is wrong, A* scores 2.71 versus 1.63 for CoT. The same pattern holds on LogicQA (3.00 vs. 1.86 for CoT-correct cases, and 2.77 vs. 1.88 for A*-correct cases). This suggests that disagreement cases are driven mainly by *coverage of the required reasoning path*, not by superficial stylistic differences.

Dataset	# disagreement cases	Annotation sample
StrategyQA	587	100
LogicQA	332	100

Table 9: Size of the disagreement pool used for human-style analysis. We consider only instances where CoT and A* produce different final predictions. In single-label QA, the *both-right* bucket is therefore expected to be empty, which is confirmed in both datasets.

Dataset	Correctness group	<i>n</i>	CoT Cov.	CoT Rel.	A*
StrategyQA	both wrong	11	1.00	2.82	
	CoT right, A* wrong	51	3.00	2.57	
	CoT wrong, A* right	38	1.63	2.47	
LogicQA	both wrong	38	1.79	2.76	
	CoT right, A* wrong	36	3.00	2.81	
	CoT wrong, A* right	26	1.88	2.81	

Table 10: Human-style annotation on 100 sampled disagreement cases per dataset. “Cov.” denotes reasoning coverage (1–3), and “Rel.” denotes relevance (1–3). Coverage measures whether the reasoning substantially reaches a correct conclusion; relevance measures the proportion of steps that stay on-topic and avoid unnecessary concepts.

C FAQs

Q1. Why are gains on HotpotQA larger than StrategyQA? HotpotQA explicitly rewards multi-hop evidence aggregation; StrategyQA is often solvable linearly.

Q2. Why include LogicQA where A* underperforms? LogicQA stresses formal sequential reasoning where branching can interfere, giving a more honest picture.

Q3. How reliable are self-reported confidence scores? They are pragmatic ranking signals (not calibrated probabilities) biasing search toward shorter, more confident traces.

Q4. How should readers interpret the admissibility result? The formal proof (Appendix J) applies only to the auxiliary depth-coverage heuristic under explicit assumptions. The main experimental heuristic is empirically validated (Appendix Section A) and also formally admissible.

Q5. Would ensemble voting capture the same complementarity? Only partially. Self-consistency recovers diversity but does not guide exploration toward promising

reasoning states.

Q6. Is goal detection based on answer patterns brittle? Somewhat. Pattern matching is fast and reproducible but can miss valid conclusions phrased differently.

D Sensitivity Analysis

All experiments are performed on a fixed stratified 300-question HotpotQA sample (seed=42). The results for budget sweep and branching factor sweep are depicted in Table 11 and Table 12.

Budget B	5	10	15	20	30
Accuracy (%)	73.0	77.3	78.0	78.7	76.0

Table 11: Budget sweep. Clear monotonic improvement $B=5 \rightarrow 20$ (+5.7 pp), diminishing returns at $B=30$.

Branching K	1	2	3	4
Accuracy (%)	77.7	78.3	76.0	80.3

Table 12: Branching factor sweep. $K=1$ (linear chain) degrades gracefully; $K=4$ achieves the highest accuracy.

E Compute Analysis

Compute-accuracy tradeoff analysis is done in Table 13.

Method	Acc.%	Tok/q	Notes
CoT	76.80	~512	single call
SC (5 traces)	78.21	~2,560	5× CoT
A*, $K=1$, $B=30$	77.7	3,372	linear chain
A*, $K=3$, $B=15$	78.0	12,452	paper default
A*, $K=3$, $B=20$	78.7	12,391	peak accuracy

Table 13: Compute-accuracy tradeoff. Token counts are measured from ablation experiments, not estimated.

F Leave-One-Benchmark-Out Transfer

To test whether the topology router generalizes across datasets (not just within-dataset), we train on all benchmarks except one and evaluate on the held-out benchmark. Results (Table 14) show that cross-dataset transfer is non-trivial—the router achieves above-chance

accuracy on most held-out benchmarks, with strongest transfer to CodeContests (79.7% on disagreement cases). However, transfer accuracy is generally lower than within-dataset CV, confirming that some dataset-specific signal remains. This motivates the feature ablation in Table 6, which shows that removing format/length features *improves* performance, suggesting the residual signal is reasoning-structural rather than format-based.

Held-Out Benchmark	Train Size	Test Size	Transfer Acc
StrategyQA	638	115	0.548
HotpotQA	519	234	0.299
LogicQA	621	132	0.462
HumanEval	722	31	0.484
MBPP	660	93	0.710
CodeContests	605	148	0.797

Table 14: Leave-one-benchmark-out transfer. The router is trained on all other benchmarks and tested on the held-out one (disagreement cases only). CodeContests and MBPP show strong transfer; HotpotQA shows weak transfer, likely because its multi-hop structure differs from other benchmarks.

G A* Component Ablation

We ablate individual components of the A* implementation on StrategyQA (500 instances, LLaMA-3.1-8B):

Configuration	Accuracy (%)
Full A* (depth + coverage heuristic)	72.4
Coverage-only heuristic (no depth)	71.0
Depth-only heuristic (no coverage)	70.0
BFS ($h = 0$, no heuristic)	69.2
CoT baseline	75.2

Table 15: A* heuristic ablation on StrategyQA. Each component contributes: full heuristic (+3.2pp over BFS), coverage adds +1.4pp over depth-only. Note: on this subset A* underperforms CoT; A* outperforms CoT on the full 2290-instance set (74.8% vs 64.5%) where multi-hop instances dominate.

H Cross-Dataset Topic Modelling

We apply Non-negative Matrix Factorisation (NMF; Lee and Seung, 1999) to all 3,660 paired question texts from the 14B Qwen-2.5 evaluation (StrategyQA 2,290; HotpotQA 719; LogicQA 651).

Pipeline. (1) Text extraction: question string for StrategyQA/HotpotQA, scenario context for LogicQA (since its query strings are formulaic). (2) TF-IDF: max_df=0.80, min_df=3, max_features=5,000, ngram_range=(1,2), English stop-words removed. (3) NMF: $k=8$, init=NNDSVD-a, max_iter=500, random_state=42. (4) Topic assignment via argmax. (5) Per-topic A*/CoT accuracy.

Validation. We select $k=8$ via NPMI coherence (0.301), the smallest k above 0.30 with interpretable topics.

Topic	Cov.%	Str.QA (A*/CoT)	Hot.QA (A*/CoT)	W
everyday_reasoning	27.7	79/79	88/73	CoT
entertainment_media	23.9	73/73	83/59	A*
historical_event	12.7	78/78	68/47	A*
geography_location	12.6	70/74	86/55	A*
general_property	9.9	73/72	63/40	A*
hypothetical_scenario	6.1	72/76	100/57	CoT
us_politics_law	4.4	72/81	91/36	A*
biography_origin	2.7	65/70	90/63	A*

Table 16: Per-topic A* vs. CoT accuracy (%) on 14B Qwen-2.5. W = overall winner across datasets.

Key finding. A* wins in 6 of 8 categories covering 66.2% of the corpus (Table 16). All A*-dominant topics share a structural property: they require resolving multi-hop entity or factual chains. The two CoT-dominant categories—*everyday_reasoning* and *hypothetical_scenario*—involve commonsense judgements where a single inference step suffices.

Representative examples. *entertainment_media* (A* +10.4 pp): “If you were on a diet, would you have to skip lunch at McDonald’s?” requires resolving menu composition plus dietary constraints across multiple facts.

historical_event (A* +2.4 pp): “Did Harry Houdini’s wife make psychics look foolish?” requires chaining Houdini’s death → wife’s debunking campaign.

geography_location (A* +9.5 pp): “Is Disney associated with Los Angeles County?” requires resolving corporate headquarters location through entity disambiguation.

everyday_reasoning (CoT +0.5 pp): “Is capturing giant squid in natural habitat impossi-

ble with no gear?” is answerable in a single commonsense step.

hypothetical_scenario (CoT +2.7 pp): “Were deaths from Apollo 13 eclipsed by other space missions?” requires a single factual comparison.

I Reproducibility Details

Models. LLaMA-3.1-8B-Instruct (main backbone), Qwen2-0.5B-Instruct (small-scale comparison), Qwen-2.5-14B-Instruct (scale experiment), DeepSeek-R1-Distill-Qwen-14B (critic/heuristic). All models served locally via Ollama (v0.3.x) or vLLM (v0.4.x) with greedy decoding for CoT and temperature sampling for A* branches (per-dataset temperatures detailed in the codebase).

Datasets. StrategyQA: full dev set, 2,290 examples. LogicQA: full test set, 651 examples. HotpotQA: stratified random sample of 500 from the fullwiki dev set (seed=42). GSM8K: test split, 1,319 examples. MATH: test split, 500 examples. HumanEval: 163 tasks. MBPP: 500 tasks. CodeContests: 500 tasks. 14B and 0.5B math experiments use the full GSM8K test split (1,319) and a 500-question MATH subset; note that the 14B A* advantage on GSM8K (+19 pp) reflects a different evaluation scale than the 8B result, which used a smaller overlapping subset in the original ACL submission.

Hardware. All experiments run on a single node with 2× NVIDIA A100 80GB GPUs. Total compute for full experimental pipeline: approximately 340 GPU-hours.

Seeds and variance. CoT uses greedy decoding (deterministic). A* uses temperature-based branching; we report single-run results but the sensitivity analysis (Appendix D) covers hyperparameter variance. Router training uses 5-fold cross-validation with stratified splits (seed=42).

Code. Full pipeline code (data preprocessing, A* search implementation, router training, evaluation scripts) is available at the anonymous repository linked in the abstract.

J Proof

We prove the result by considering three cases.

Case 1: Root node ($d_n = 0$). By Assumption 1, any valid goal trace requires at least $H_{\min}$ future transitions from the root. Since each future transition costs at least $c_{\min}$, the true remaining cost satisfies

$$h^*(n) \geq H_{\min} c_{\min}.$$

At the root, $d_n = 0$, so

$$h_{\text{depth}}(n) = \omega_d H_{\min}.$$

Also, because $E(n) \in [0, 1]$, we have

$$h_{\text{coverage}}(n) = \omega_c(1 - E(n)) \leq \omega_c.$$

Therefore,

$$h(n) = \omega_d H_{\min} + \omega_c(1 - E(n)) \leq \omega_d H_{\min} + \omega_c.$$

Using $\omega_d + \omega_c \leq c_{\min}$, and hence $\omega_d \leq c_{\min}$, we obtain

$$\begin{aligned} \omega_d H_{\min} + \omega_c &= (H_{\min} - 1)\omega_d + (\omega_d + \omega_c) \\ &\leq (H_{\min} - 1)c_{\min} + c_{\min} = H_{\min} c_{\min} \leq h^*(n) \end{aligned} \quad (7)$$

Thus $h(n) \leq h^*(n)$ at the root.

Case 2: Intermediate non-goal node ($0 < d_n < H_{\min}$). Let

$$m = H_{\min} - d_n.$$

Since $0 < d_n < H_{\min}$, we have $m \geq 1$. By Assumption 1, at least m further transitions are required before reaching a goal. Since each transition costs at least $c_{\min}$,

$$h^*(n) \geq m c_{\min}.$$

For such a node,

$$h_{\text{depth}}(n) = \omega_d m.$$

Again, since $E(n) \in [0, 1]$,

$$h_{\text{coverage}}(n) = \omega_c(1 - E(n)) \leq \omega_c.$$

Therefore,

$$h(n) = \omega_d m + \omega_c(1 - E(n)) \leq \omega_d m + \omega_c.$$

Using $\omega_d + \omega_c \leq c_{\min}$ and $\omega_d \leq c_{\min}$, we get

$$\begin{aligned} \omega_d m + \omega_c &= (m - 1)\omega_d + (\omega_d + \omega_c) \\ &\leq (m - 1)c_{\min} + c_{\min} = m c_{\min} \leq h^*(n). \end{aligned} \quad (8)$$

Thus $h(n) \leq h^*(n)$ for all intermediate non-goal nodes.

Case 3: Deep non-goal node ($d_n \geq H_{\min}$). By Assumption 2, even though the node has already reached the minimum depth, any non-goal node still requires at least one additional transition before reaching a goal. Hence,

$$h^*(n) \geq c_{\min}.$$

For $d_n \geq H_{\min}$, the depth term is zero:

$$h_{\text{depth}}(n) = \omega_d \max(0, H_{\min} - d_n) = 0.$$

Thus,

$$h(n) = \omega_c(1 - E(n)).$$

Since $E(n) \in [0, 1]$,

$$h(n) \leq \omega_c.$$

Because $\omega_c \leq \omega_d + \omega_c \leq c_{\min}$, we have

$$h(n) \leq c_{\min} \leq h^*(n).$$

Thus $h(n) \leq h^*(n)$ also holds for deep non-goal nodes.

Since the inequality $h(n) \leq h^*(n)$ holds in all three cases, the heuristic is admissible for all non-goal nodes under Assumptions 1 and 2. $\square$

Interpretation. The theorem should be read as a conservative guarantee for the depth-plus-coverage heuristic under explicit assumptions. It does not claim that the heuristic is optimal, universally best, or valid for every dataset-specific implementation. In particular, the theorem does not cover the LogicQA-specific remaining-depth/logical-complexity heuristic. Its role is narrower: it shows that the depth-plus-coverage heuristic used in StrategyQA and HotpotQA can be interpreted as a conservative lower bound on the remaining path cost when the stated assumptions and weight condition hold. The empirical robustness of the overall A* procedure is evaluated separately through sensitivity analysis over node budget and branching factor.

K Discussion

The CoT/A* boundary is structural, not dataset-specific. The strongest empirical finding is that A* advantage tracks branching coefficient across clusters (-0.4 pp in C1 to

+11.4 pp in C4, Table 3), not benchmark labels. This means the question “should I use search?” is better answered by looking at the input’s structure than by knowing which dataset it comes from.

Pre-generation routing beats trace inspection. This was the most surprising result to us. The topology router sees zero generated tokens yet outperforms critics that read both full traces. The explanation is the Fluency-Correctness Asymmetry: in 38% of disagreements, the wrong answer comes wrapped in the more fluent trace, fooling any judge that conflates readability with correctness. The implication is that adaptive inference should start with structural triage, not trace evaluation. The critic-based heuristic itself operates conservatively (mean 0.43 vs. human 0.58, $\epsilon_{\max} = 0.14$), providing practical confidence without formal admissibility.

Substantial room remains. The topology router recovers only 26.1% of the oracle gap. The remaining 74% likely involves factors we do not yet model: whether the base LLM has the right parametric knowledge for a given question, whether the search depth is matched to the problem, and how model capacity interacts with topology (the 14B results in Table 1 hint at complex interactions).

L Code Generation Analysis

This section provides detailed analysis of the code-generation benchmarks (HumanEval, MBPP, CodeContests) that complement the main results in Table 1.

Topology router on code. The topology router recovers 67.8% of the CoT–A* oracle gap on HumanEval at 8B (CV accuracy: 0.867 on 29 disagreement cases, using code-adapted structural features). On MBPP, where prompts average only 14 words with near-zero branching complexity, the router achieves −13.9% gap recovery (worse than always-CoT). This contrast confirms the topology thesis: routing succeeds when input structure provides discriminative signal, and correctly abstains when features are degenerate. The top predictive features on HumanEval are prompt length (0.757 importance), negation count (0.089), and depth estimate (0.055).

Multi-return tasks: the code analogue of branching topology. On HumanEval, tasks requiring multi-value returns (tuples, composite outputs; $n=11$) exhibit a striking scale-dependent interaction (Table 17). At 8B, CoT achieves 100% while A* achieves 81.8% (−18.2 pp). At 14B, the pattern reverses: A* achieves 100% while CoT drops to 81.8% (+18.2 pp). This “capability threshold” effect suggests that search on structurally complex tasks requires sufficient base-model competence to generate meaningful alternatives—the code analogue of the branching-coefficient finding in QA.

Scale	Multi-return CoT	Multi-return A*	A* advantage
0.5B	45.5%	18.2%	−27.3 pp
8B	100.0%	81.8%	−18.2 pp
14B	81.8%	100.0%	+18.2 pp

Table 17: Multi-return tasks on HumanEval ($n=11$) across scales. A* achieves perfect accuracy only at 14B, where the base model is capable enough to generate meaningful alternative solutions.

Category analysis (MBPP). Table 18 breaks down MBPP performance by problem type. A* shows consistent advantages on string manipulation (+4.3 pp at 8B) and logic/validation tasks (+5.0 pp at 14B), while CoT leads on list operations (−2.4 to −3.3 pp) and math (−7.5 to −1.3 pp). This maps onto the QA pattern: A* helps when multiple conditions must be verified, while CoT is preferable for tasks with a clear algorithmic transformation.

Category	n	8B (A* adv.)	14B (A* adv.)
String	93	+4.3 pp	+1.1 pp
Logic	60	−1.7 pp	+5.0 pp
Other	55	+3.6 pp	+1.8 pp
List	246	−2.4 pp	−3.3 pp
Math	227	−7.5 pp	−1.3 pp

Table 18: MBPP A* advantage (pp over CoT) by problem category. A* helps on string and logic tasks (multi-condition checking), while CoT leads on list and math tasks (clear algorithmic transformations).

CodeContests: search helps across scales. CodeContests exhibits strong A* advantage at multiple scales. At 8B, A* outper-

forms CoT by +9.6 pp (87.0% vs. 77.4%), with 94 instances (18.8%) solved exclusively by A* and an oracle of 96.2%. The topology router reaches 95.0%—recovering 87% of the oracle gap, the strongest single-dataset routing result in this work. At 0.5B, the advantage is even larger: A* gives +18.4 pp over CoT (68.2% vs. 49.8%), with 26.2% solved exclusively by A*. This pattern—search helping on algorithmic tasks where the model can generate meaningful alternatives—mirrors HotpotQA and confirms that CodeContests problems have genuine branching topology despite their short specifications.

Compute. A* uses only 1.7–3.4× the time of CoT on code tasks (vs. ~24× on QA), because code solutions are shorter and require fewer search nodes. This makes topology-based routing particularly practical in the code domain.

Qualitative examples: A* success on code.

HumanEval/1 — A* Correct, CoT Wrong

Task: `separate_paren_groups`: Given a string with multiple nested parenthesis groups, separate them into individual balanced strings.

CoT (×): Generates a solution that fails to properly handle space removal before parsing, leading to incorrect group boundaries.

A* (✓): First removes all spaces, then uses a stack-based approach to track nesting depth and correctly identifies group boundaries. The search explores multiple parsing strategies before committing to the stack-based solution.

HumanEval/33 — A* Correct, CoT Wrong

Task: `sort_third`: Return a list identical to the input except that values at indices divisible by three are sorted.

CoT (×): Misinterprets the indexing condition, sorting elements at wrong positions.

A* (✓): Correctly extracts elements at indices 0, 3, 6, ..., sorts them independently, and reinserts them. The multi-constraint nature (preserve some positions, sort others) benefits from search exploring different index-selection strategies.

Qualitative examples: CoT success on code (A* overthinks).

HumanEval/0 — CoT Correct, A* Wrong

Task: `has_close_elements`: Check if any two numbers in a list are closer than a threshold.

CoT (✓): Directly implements the $O(n^2)$ pairwise comparison—simple, correct, and complete.

A* (×): Attempts an optimization by sorting first, but introduces an incorrect compound condition (`abs(numbers[i] - numbers[i-1]) < threshold and abs(numbers[i] - numbers[i+1]) < ...`), adding an unnecessary constraint that breaks correctness.

HumanEval/14 — CoT Correct, A* Wrong

Task: `all_prefixes`: Return all prefixes of a string from shortest to longest.

CoT (✓): One-line list comprehension: `[string[:i+1] for i in range(len(string))]`—direct and correct.

A* (×): Generates `range(len(string) + 1)`, which includes an empty string as the first prefix, failing the test cases. Search explored an off-by-one variant that a single linear pass would not have considered.

Pattern. These examples confirm the same pattern as QA: A* helps when multiple constraints must be simultaneously satisfied (group boundaries + nesting depth in HumanEval/1; index selection + sorting in HumanEval/33). CoT remains preferable for di-

1465
1466
1467

1468

1469

1470

1471

rect algorithmic implementations where the solution path is obvious and exploration risks introducing unnecessary complexity.

L.1 Additional HumanEval Examples

HumanEval/43 — A* Correct, CoT Wrong

Task: pairs_sum_to_zero: Return True if there are two distinct elements in the list that sum to zero.

CoT (×): Generates a solution that checks `l[i] + l[j] == 0` but uses incorrect index bounds or fails to enforce distinctness of positions.

A* (✓): Explores a set-based approach—for each element, checks if its negation exists in a seen-set, correctly handling the distinct-element constraint.

HumanEval/67 — A* Correct, CoT Wrong

Task: fruit_distribution: Given a string like “5 apples and 6 oranges” and total n , return the count of other fruits (n minus apples minus oranges).

CoT (×): Fails to correctly parse the string—misidentifies which numbers correspond to which fruits.

A* (✓): The search explores different parsing strategies and settles on extracting all integers from the string, then computing $n - \text{sum}(\text{extracted})$.

HumanEval/57 — CoT Correct, A* Wrong

Task: monotonic: Return True if list elements are monotonically increasing or decreasing.

CoT (✓): Directly checks `l == sorted(l)` or `l == sorted(l, reverse=True)`—simple and correct.

A* (×): Explores pairwise comparison approaches but introduces an off-by-one error in the iteration, failing on edge cases.

L.2 MBPP Examples (LLaMA-3.1-8B)

1472

MBPP/118 — A* Correct, CoT Wrong

Task: Convert a string to a list (expected: split on whitespace).

CoT (×): Generates list(s) which splits the string into individual characters, not words.

A* (✓): Explores multiple splitting strategies and correctly identifies whitespace as the delimiter, generating `input_string.split()`.

1473

MBPP/130 — A* Correct, CoT Wrong

Task: Find the item with maximum frequency in a given list.

CoT (×): Uses Counter but returns the count instead of the item, or fails to handle the most-common extraction correctly.

A* (✓): Builds a frequency dictionary manually and correctly returns the key (not value) with highest count.

1474

MBPP/100 — A* Correct, CoT Wrong

Task: Find the next smallest palindrome of a specified number.

CoT (×): Increments and checks palindrome but fails on edge cases (e.g., numbers like 999 that require length change).

A* (✓): Explores both increment-and-check and mirror-based approaches, handling the length-boundary edge case.

1475

MBPP/14 — CoT Correct, A* Wrong

Task: Find the volume of a triangular prism.

CoT (✓): Directly applies the formula $V = \frac{1}{2} \times \text{base} \times \text{height} \times \text{length}$ with the correct function signature.

1476

A* (×): Generates a function with wrong parameter names (base_area, height instead of three separate parameters), producing incorrect results.

MBPP/30 — CoT Correct, A* Wrong

Task: Count all substrings starting and ending with the same character.

CoT (✓): Simple nested loop checking $s[i] == s[j]$ for all valid (i, j) pairs—correct and complete.

A* (×): Attempts an optimization using character frequency counts but miscounts single-character substrings.

MBPP/32 — CoT Correct, A* Wrong

Task: Find the largest prime factor of a given number.

CoT (✓): Standard trial division from 2 upward, keeping track of the largest factor—correct.

A* (×): Overcomplicates with a separate primality test function and fails to correctly iterate through all factors.

L.3 CodeContests Examples (Qwen-0.5B)

CodeContests examples are drawn from the 0.5B model, where A* provides the largest advantage (+18.4 pp). At 8B and 14B, CoT solves nearly all tasks directly, leaving no A*-only instances for qualitative analysis.

CC/0 — A* Correct, CoT Wrong (0.5B)

Task: Return the length of the longest strictly increasing subsequence.

CoT (×): Generates a skeleton with docstring but fails to produce a working implementation (empty or trivially wrong body).

A* (✓): Explores multiple approaches and produces a correct dynamic programming solution with $O(n^2)$ complex-

ity.

CC/4 — A* Correct, CoT Wrong (0.5B)

Task: Count all solutions to the N-Queens problem.

CoT (×): At 0.5B scale, CoT generates an empty or syntactically invalid function body—the model cannot produce backtracking logic in a single linear pass.

A* (✓): Search explores candidate board placements step-by-step, assembling a correct backtracking solution with row-by-row queen placement and constraint checking.

CC/6 — A* Correct, CoT Wrong (0.5B)

Task: Implement the trapping rain water algorithm.

CoT (×): Generates a function with correct docstring but an incomplete or incorrect body that does not compute trapped water.

A* (✓): Explores precomputation of left-max and right-max arrays, correctly implementing the $O(n)$ two-pass approach.

CC/7 — CoT Correct, A* Wrong (0.5B)

Task: Find all valid combinations that sum to a target.

CoT (✓): Generates a correct recursive backtracking solution with proper base cases and candidate enumeration.

A* (×): The search branches lead to conflicting partial implementations that fail to merge into a coherent recursive structure.

CC/14 — CoT Correct, A* Wrong (0.5B)

Task: Count solutions to the N-Queens problem (variant).

CoT (✓): Produces a complete backtracking solution with column, diagonal, and anti-diagonal tracking.

A* (×): Generates a structurally similar but incorrect solution—the column-tracking logic has an off-by-one error introduced during branch selection.

CC/24 — CoT Correct, A* Wrong (0.5B)

Task: Count solutions to N-Queens.

CoT (✓): Another variant where CoT’s single-pass generation produces correct queen placement logic.

A* (×): Search explores redundant branches for the same subproblem, wasting budget and ultimately selecting a suboptimal partial solution as the final answer.

Cross-dataset pattern. Across all three code benchmarks, A* succeeds when the problem requires *exploring alternative solution strategies* (different data structures for LIS, different parsing approaches for string problems, multiple constraint-satisfaction strategies for N-Queens). CoT succeeds when the solution follows a *well-known algorithmic template* (backtracking, trial division, direct formula application) that can be generated correctly in a single forward pass. This mirrors the QA finding: A* helps on branching problems; CoT helps on linear ones.

LLM-as-Judge

Purpose. Compare two candidate reasoning trajectories and decide which one solves the query more *faithfully* and *correctly*.

Inputs.

- Context
- Query
- Path A reasoning trace
- Path B reasoning trace

Evaluation criteria.

- **Faithfulness:** uses only the provided context and avoids invented constraints
- **Logical validity:** each step follows causally and coherently from prior steps
- **Grounded conclusion:** the final answer directly resolves the query from the established reasoning

Required output. The model first reasons internally about both traces, then returns only the winning path in constrained XML:

`<winner>A</winner>` or
`<winner>B</winner>`

LLM-as-Critic

Purpose. Diagnose the weaknesses of two reasoning traces and determine which trace is better overall.

Inputs.

- Question
- Answer options
- Trace A with predicted answer
- Trace B with predicted answer

Evaluation criteria.

- whether the reasoning actually entails the predicted answer
- whether there are logical fallacies or unsupported jumps
- whether any claims are hallucinated or ungrounded

Required output. The model returns a structured JSON critique with flaws for both traces and the preferred trace:

```
{
  "trace_A_flaws": "...",
  "trace_B_flaws": "...",
  "better_trace": "A/B"
}
```

Figure 3: LLM-based routing prompts. The *judge* prompt is verdict-oriented: it compares two reasoning paths under faithfulness and logical soundness criteria and outputs only the winner. The *critic* prompt is diagnosis-oriented: it identifies flaws in each trace before selecting the better one in structured JSON form.